

Parallel Programming Homework 2

Mandelbrot Set

CS 107062546 Chung-Kai Yang 楊仲愷

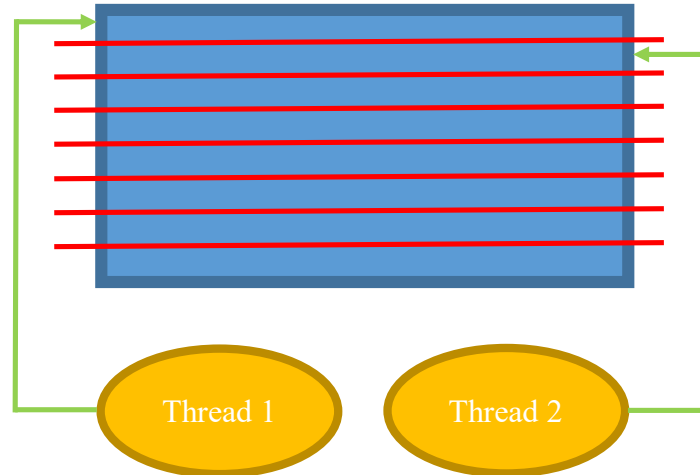
1. Implementation

(1) How you implement each of requested versions, especially for the hybrid parallelism.


a. Pthread

As we can see in the implementation of Mandelbrot Set, I take advantage of dynamically calculate height (i.e., divide with row) with pthread.

Step 1:

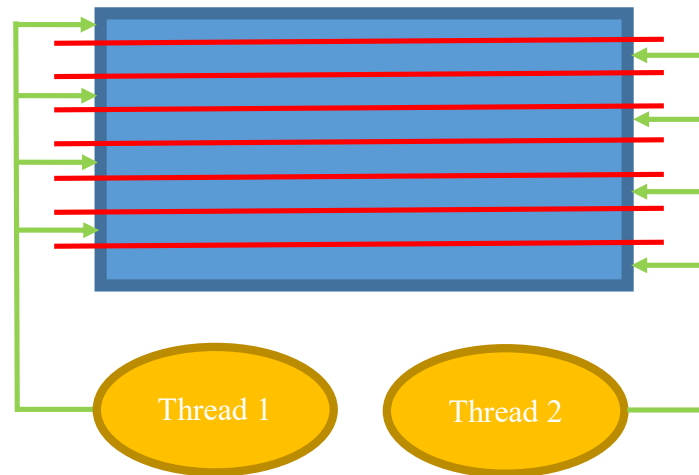


Step 2 (when thread 2 finish):



b. Hybrid (MPI + OpenMP)

In this section, I utilize the same concept to fetch the result of Mandelbrot Set. Here, I have tried static (i.e., distribute fixed chunk size before executing) and dynamic (i.e., when each thread finish, get the next one with round robin scheduling strategy) version.



(2) *How do you partition the task?*

a. Pthread

In this version, as I describe before, I divide the image by rows. And each thread will get one to compute; when each one finish, it will get the next one until all the rows are done.

b. Hybrid (MPI + OpenMP)

The partition in this section is a little different. I distribute the image into rows, and each rank get a part of them (e.g., the height is 7, and the rank size is 3; thus, rank 0 will get height 0, 3, 6; rank 1 will get 1, 4, 7; rank 2 will get 2, 5.) because of the probability for the amount of computation.

(3) *What technique do you use to reduce execution time and increase scalability?*

a. Pthread

In this part, I have tried some other versions, like partition image into pixels, or into columns. However, it seems partition into rows is the best choice. Furthermore, I also have tried partition statically, but it doesn't perform well. Therefore, I partition image into rows and dynamically obtain the next one to compute to increase scalability.

b. Hybrid (MPI + OpenMP)

In this part, I have tried some other versions as well, like partition image into pixels or into columns. However, it seems partition into rows is the best choice. Furthermore, I have tried some improvements to accelerate. First, I change the for loop step into rank size. Second, I replace and decrease the number of branches. And this performs well.

(4) *Other efforts you made in your program.*

As mentioned before, I have tried some other versions as well, like partition image into pixels or into columns. Also, I change the for loop step into rank size. And I replace and decrease the number of branches, etc.

2. Experiment & Analysis

i. Methodology

(a) System Spec

Run on cluster.


(b) Performance Metrics

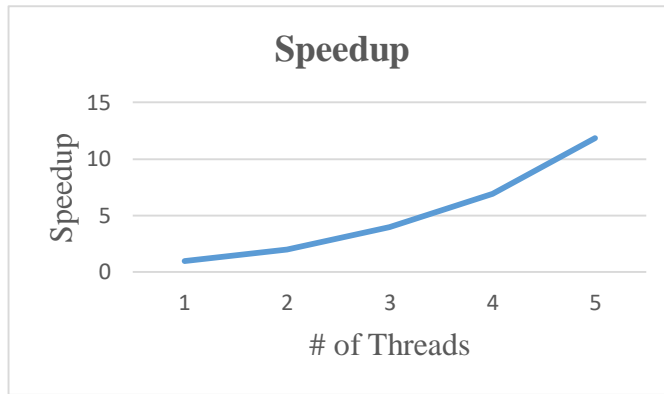
(1) Computing time: Time for accessing CPU. (e.g., calculate Mandelbrot Set)

ii. Plots: Scalability & Load Balancing

(a) Pthread



	CPU	Total Time	Speedup
1	119.682365	119.6824	1
2	59.865865	59.86587	1.999175
4	30	29.98975	3.990775
8	17.245502	17.2455	6.939918
12	10.122554	10.12255	11.82334



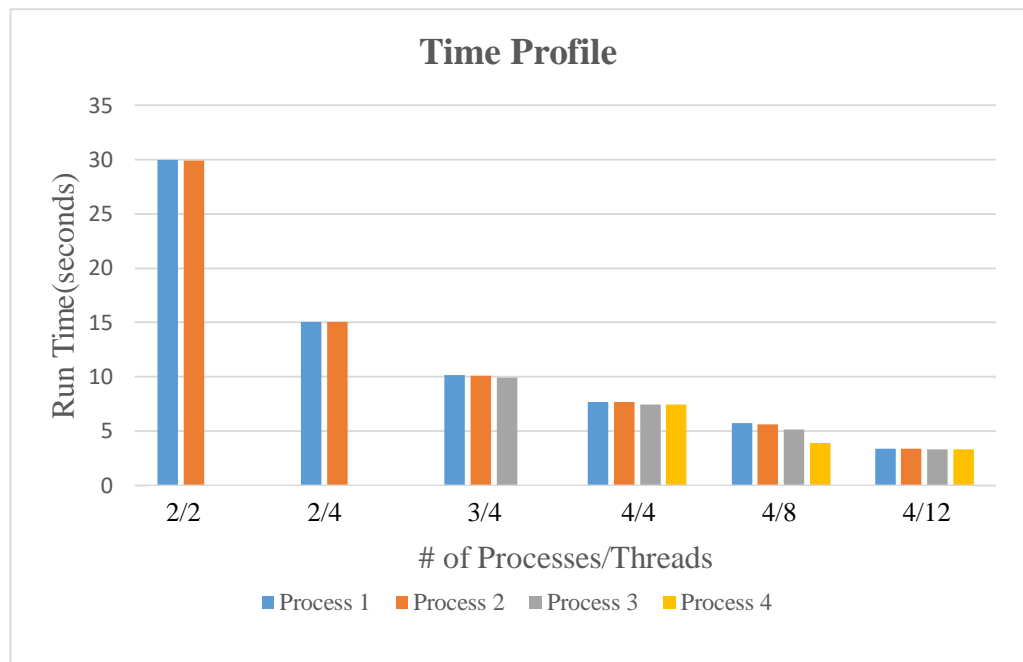


(b) Hybrid (MPI + OpenMP)


Process	Thread	CPU	Total Time	Speedup
1	1	119.6905	119.6905	1
1	2	59.97039	59.97039	1.995826
2	2	30	29.98946	3.991086
2	4	15.04867	15.04867	7.95356
3	4	10.13889	10.13889	11.80509
4	4	7.690905	7.690905	15.5626
4	8	5.750069	5.750069	20.81549
4	12	3.360959	3.360959	35.61201




(c) Load Balance



(d) Other Experiment





3. Discussion

(a) Scalability:

- (1) **Strong Scaling:** No matter Pthread or Hybrid version, according to figures mentioned before, as the number of processes increasing, the total amount of computation holds; therefore, the calculation quantity each processor has to process will decrease, and the computing time will also lower with parallel computing.
- (2) **Weak Scaling:** I take the example below (i.e., case *slow04*), and I test different image size (i.e., 1920×1080 , 960×540). First, I lower 2 times than the height and width before; then, the total computing quantity 4 times, and the time it takes is almost the same.

Pthread:

```
[pp19s33@apollo31 hw2]$ srun -n1 -c12 ./hw2a out.png 17326507 -0.5506164691618783 -0.5506164628264113 0.62734454371181
131 0.6273445403522527 1920 1080
12 cpus available
10.122919
[pp19s33@apollo31 hw2]$ srun -n1 -c3 ./hw2a out.png 17326507 -0.5506164691618783 -0.5506164628264113 0.62734454371181
31 0.6273445403522527 960 540
3 cpus available
10.080187
```

Hybrid:

```
[pp19s33@apollo31 hw2]$ srun -n4 -c6 ./hw2b out.png 17326507 -0.5506164691618783 -0.5506164628264113 0.62734454371181
31 0.6273445403522527 1920 1080
4.972213
[pp19s33@apollo31 hw2]$ srun -n2 -c3 ./hw2b out.png 17326507 -0.5506164691618783 -0.5506164628264113 0.62734454371181
31 0.6273445403522527 960 540
5.023840
```

(b) Load Balance

- (1) **Pthread:** As I mentioned before, I distribute every thread to calculate each row. Then, when each thread finish, it will fetch the next, and it is like a dynamic version.

(2) **Hybrid:** Because the computation quantity of each row is not uniform, I take the same concept as pthread version, I divide the image by rows. And each rank will only get a part of them (e.g., the height is 7, and the rank size is 3; thus, rank 0 will get height 0, 3, 6; rank 1 will get 1, 4, 7; rank 2 will get 2, 5.).

(c) Other Experiment

From the first figure, it shows the dynamic version is more efficient than static version, and I think maybe static version doesn't distribute uniformly enough to cause the result

From the second figure, it shows if I divide by pixels it may not lower the time because when a thread finish, it fetches the next pixel to calculate; however, it happens too frequently to lock the pixel and other threads are always waiting for release the mutex.

4. Conclusion

From this homework, I am getting more familiar with pthread and hybrid (MPI + OpenMP) version. Therefore, according to every result, I try different ways to improve the latency. And I also try dynamic and static or chunk size. Also, I learn the importance of load balance. In spite of the fact that I take advantage of parallel computing, I still can't lower the time.